GURPREET KAUR

Aspiring Data Scientist

gunjan010908@gmail.com | +91 8130863925 | linkedin.com/in/gurpreet-kaur-55100341a | github.com/gunjan01-source | gunjan01-source.github.io/my-portfolio

PROFILE SUMMARY

Detail-oriented BCA graduate (2025) with a strong foundation in data science and machine learning. Skilled in Python, data manipulation, visualization, and BI tools (Power BI, Tableau), with hands-on project experience turning raw data into actionable business insights and building predictive models. Eager to apply analytical and modeling skills to an entry-level Data Scientist role.

TECHNICAL SKILLS

Programming: Python (Pandas, NumPy, Scikit-learn), SQL

Data Visualization: Power BI, Tableau, Matplotlib

Databases: MySQL, MongoDB

Concepts: Data Cleaning, EDA, Regression, Machine Learning Fundamentals

PROJECTS

Retail Sales Performance Analysis (Superstore)

Python, Pandas, NumPy, Matplotlib, MySQL, Power BI | github.com/gunjan01-source/superstore-analysis

- Cleaned and analyzed 9,900+ retail orders using Python and Pandas, then loaded the dataset into MySQL and wrote queries to surface category, region, and segment-level KPIs.
- Built a 2-page interactive Power BI dashboard with KPI cards, a conditionally-formatted Segment × Category profit matrix, region/segment slicers, a Discount-vs-Profit scatter plot, and an AI Key Influencers visual to diagnose loss-making orders.
- Found Technology generated the highest sales and the West region drove the highest revenue, while high-discount orders (40%+) consistently turned unprofitable across categories.
- Identified Furniture as a systemic loss driver, with negative profit across every customer segment despite solid sales volume, and documented findings in a written insights report.

House Price Prediction

Python, Pandas, Scikit-learn, Matplotlib | github.com/gunjan01-source/house-price-prediction

- Built a regression pipeline to predict house prices from structural and location features (area, bedrooms, bathrooms, age, location, condition), including data cleaning and feature engineering (total rooms, sqft-per-room, age buckets).
- Trained and compared Linear Regression, Random Forest, and Gradient Boosting models using a scikit-learn ColumnTransformer pipeline for scaling and one-hot encoding.
- Selected Linear Regression as the best-performing model (R^2 of 0.96, RMSE ≈ $58K), and visualized feature coefficients and predicted-vs-actual performance to validate the fit.

EDUCATION

Bachelor of Computer Applications (BCA)

Jamia Hamdard University | Batch of 2025